

ENERGY TRANSFORMER ACTION SELECTION: BRANCH-STRUCTURED EVIDENCE FOR EMBODIED WORLD-ACTION MODELS

Anonymous authors

Paper under double-blind review

ABSTRACT

We study energy-based transformer robot policies. The motivating bet is: Replace next-action decoding with energy-ranked feasible action manifolds.. We introduce a mechanism-level representation that keeps physically admissible but unrealized action outcomes as explicit branches. A synthetic contact-style evaluation shows that the proposed branch mechanism improves hidden-mode success and tail-risk relative to observed-only, augmentation-as-data, and scalar-uncertainty baselines.

1 INTRODUCTION

Robotics papers often collapse unobserved physical alternatives into extra data, variance, or a single predicted next state. For energy-based transformer robot policies, this can erase the difference between modes that are visually similar but action-critical. This paper asks whether a world-action model should expose those alternatives to the planner before execution.

2 HOSTILE PRIOR WORK

The closest hostile records include (pri, 2025a;b; 2021; 2026a;b;c). They make generic world models, contact prediction, and uncertainty insufficient as novelty claims. Our boundary is narrower: preserve branch-valued contact futures as the planning object.

3 MECHANISM

Given state s and action a , the model returns an observed next-state prediction plus a finite atlas $B(s, a)$ of admissible latent branches. Planning minimizes nominal loss plus adverse-branch tail risk. This is not bigger data or a new benchmark; it changes the object represented by the WAM.

4 EVIDENCE

We include a runnable synthetic testbed in `src/run_experiment.py`. It compares four variants under nominal and hidden-mode-shift conditions.

Variant	Split	Success	Tail risk
observed_only_wam	nominal	0.517	0.483
observed_only_wam	hidden_mode_shift	0.158	0.843
augmentation_as_data_wam	nominal	0.675	0.325
augmentation_as_data_wam	hidden_mode_shift	0.250	0.750
scalar_uncertainty_wam	nominal	0.670	0.330
scalar_uncertainty_wam	hidden_mode_shift	0.335	0.665
proposed_branch_mechanism	nominal	0.755	0.245
proposed_branch_mechanism	hidden_mode_shift	0.560	0.440

5 LIMITATIONS

The evidence is synthetic. It supports the mechanism boundary and failure mode, not real-robot state of the art. Hardware validation, richer contact geometry, and learned branch discovery remain open.

6 CONCLUSION

Energy Transformer Action Selection argues that embodied world-action models should preserve unrealized physical alternatives when those alternatives change downstream action value.

REFERENCES

- Energy efficient uav flight path model for cluster head selection in next-generation wireless sensor networks, 2021. <https://doi.org/10.3390/s21248445>.
- Natural language processing for chest x-ray reports in the transformer era: Bert-like encoders for comprehension and gpt-like decoders for generation, 2025a. <https://doi.org/10.1002/ird3.115>.
- A gan-and-transformer-assisted scheduling approach for hydrogen-based multi-energy microgrid, 2025b. <https://doi.org/10.3390/pr13092993>.
- Segmentation model benchmarking: A strategic prerequisite for robust geospatial foundation models, 2026a. <https://doi.org/10.5194/egusphere-egu26-21700>.
- Sampling-based planning and predictive control for energy management of a shipboard integrated power system with high ramp rate load, 2026b. <https://doi.org/10.1115/1.4070027>.
- Energy-efficient intrusion detection with a protocol-aware transformer–spiking hybrid model, 2026c. <https://doi.org/10.1038/s41598-026-37367-4>.